

# Spreadsheets & Hallucinations

A PR coverage data manifesto

# Introduction

For fifteen years we've been thinking exclusively about coverage data.

We've seen countless reports, dashboards, PDFs, and monthly summaries built from spreadsheets that classify coverage in many different ways. Despite the variety of organisations, from global brands to major agencies to small in-house teams, the same problems repeat everywhere.

Most teams are trying to extract insight from systems that were never designed to support it. Coverage is classified differently in every organisation, often rebuilt from scratch, and stored in fragile documents that don't scale, don't collaborate well, and don't survive long enough to become genuinely useful.

For years, this was frustrating but manageable. Reporting was occasional and retrospective. It was a way of justifying what had already happened.

In the last 12 to 18 months, that changed. Organisations now expect PR to deliver on-demand insight, not just periodic summaries. They want market intelligence, narrative signals, and actionable answers. They want it quickly.

Waiting in line to pay for a specialist analysis, or working backwards through broken spreadsheets no longer fits the pace at which decisions are made. Insight that arrives next month is often already too late.

At the same time, the rapid rise of AI has raised the stakes. Mathematical machines are now being pointed at PR data and asked to surface patterns, insights, and answers. When that data is unstructured, inconsistent, or poorly defined, the results are misleading. Confident, plausible-sounding conclusions get built on weak foundations. This is a structural problem.

This report exists because we believe the PR industry has reached a moment where it needs to do some foundational work. Not another framework. Not another metric. A way of organising coverage data so it can be trusted, queried, and learned from by people and by machines.

We are building our products around this standard because we believe it is the right foundation. But the ideas in this report are broader than any one tool. Whether you use CoverageBook, a spreadsheet, or something else entirely, the underlying principles still apply. If PR is to move from justification to intelligence, it needs structure. If AI is to help rather than mislead, it needs clean inputs.

If the industry is to keep up with modern decision-making, it needs answers now, not later.

This report is a contribution towards that work.

# Part 1 - The state of PR measurement today

To understand where PR measurement needs to go next, it helps to be clear about where it has come from.

Ten years ago, PR measurement was largely focused on outputs. Success was reported in terms of volume: how many pieces of coverage were secured, where they appeared, and how much notional reach they generated.

Reports were typically delivered as PDFs at the end of a campaign or month, summarising clips, publication logos, and headline numbers.

AVEs, despite being widely criticised, were still commonly used as a shorthand for value. The aim was reassurance rather than insight, evidence that activity had taken place and visibility had been achieved.

That approach reflected both the tools available at the time and the expectations placed on PR. Measurement was retrospective, static, and largely disconnected from the rest of the organisation's data.

Over time, that began to change.

As marketing, analytics, and digital channels matured, expectations of PR evolved with them. Organisations started asking harder questions. Coverage volume alone was no longer enough. Attention shifted towards outcomes: narrative penetration, reputation, leadership visibility, message pull-through, and the role PR played in influencing behaviour.

PR reporting began to move closer to marketing measurement. Teams started looking for correlations with branded search, website traffic, social engagement, recruitment interest, and brand tracking. The conversation shifted from what coverage did we get? to what did that coverage actually do?

This was a necessary and healthy evolution. But even as expectations increased, the underlying systems PR relied on did not fundamentally change.

Much of the industry was still working with unstructured data, manually classified in spreadsheets and static reports. Insight was possible, but slow. Analysis was often bespoke and resource-intensive. And reporting remained largely periodic, useful for reflection, less useful for decision-making in real time.

In the last year, the ground shifted again.

AI has rapidly become the interface through which people access information. Search is no longer just a list of links; it is increasingly a synthesised answer. Tools like ChatGPT, Perplexity, and others don't simply retrieve content, they interpret it, summarise it, and present it as knowledge.

This has two profound implications for PR.

First, PR content has become training data.

Journalism, expert commentary, brand narratives, and thought leadership are now part of the raw material AI systems draw from to explain the world. PR doesn't just shape human understanding anymore; it shapes machine understanding, and those machines, in turn, shape how people form opinions and make decisions.

Second, the expectations placed on PR reporting have accelerated dramatically.

Organisations no longer want occasional justification of past activity. They want on-demand insight. They expect PR to contribute to market intelligence: to surface emerging narratives, track leadership visibility, identify reputational risk, and inform decisions as they are being made.

Insight that arrives weeks later is often already out of date. Waiting in a queue for a data science team, or manually pulling analysis from fragmented spreadsheets, no longer fits the pace at which businesses operate.

PR is being asked to move from coverage reporting to coverage intelligence.

This is where the problem becomes clear.

AI is exceptionally good at spotting patterns, summarising information, and generating plausible-sounding answers. But it has a critical weakness: it is only as reliable as the data it is given.

When AI systems are pointed at unstructured PR data, inconsistent tags, unclear categories, fragmented datasets, they fail confidently. They produce insights that look coherent, but rest on shaky foundations. Hallucinations, false correlations, and misleading conclusions are not edge cases; they are predictable outcomes of poor structure.

What was once an inefficiency has now become a risk.

Unstructured PR data doesn't just limit insight, it actively undermines trust in the answers being produced. It creates false confidence at exactly the moment organisations are relying more heavily on automated analysis. This is the paradox PR now faces.

PR has never been more important. Its influence extends beyond audiences to the systems that mediate knowledge itself. At the same time, the way PR data is collected, organised, and analysed has not kept pace with that responsibility.

The challenge is not a lack of activity, ambition, or expertise within PR teams. It is the absence of a shared structural foundation.

If PR is now feeding AI systems, if it is expected to deliver real-time intelligence rather than retrospective reports, then the way coverage data is organised suddenly matters a great deal.

That brings us to the question at the heart of this report.

If insight depends on structure, what happens when an entire industry has never agreed on how to structure its data?

## Part 2 - The spreadsheet nightmare

Across the PR industry, there is one tool that appears almost everywhere. It doesn't matter whether the team sits inside a global agency, a multinational brand, or a small in-house function. It doesn't matter how sophisticated their strategy is, or how experienced the practitioners are. At some point, nearly all PR teams end up classifying their coverage in a spreadsheet.

Thousands of teams. Hundreds of thousands of practitioners. All trying to organise and analyse coverage at scale using tools that were never built for that purpose. This isn't a failure of effort or intent. It's a symptom of an industry that has never been given a better alternative.

Spreadsheets are powerful tools for calculation. They are poor tools for classification, collaboration, and long-term analysis. In PR, the same problems show up repeatedly.

They struggle with collaboration. As soon as multiple people are involved, spreadsheets become fragile. Cells are overwritten, dropdowns are edited, formulas break, and version control becomes a guessing game.

They lack durability. A carefully constructed sheet works until someone new joins, a campaign changes, or leadership asks a new question the structure wasn't designed to answer.

They encourage inconsistency. Column names drift. Categories multiply. Labels mean slightly different things to different people. What started as a neat system degrades into something no one trusts.

And critically, spreadsheets don't scale. They don't accumulate value over time. Each new reporting period often means copying, pasting, and adapting rather than building on a stable foundation. By the time insight is needed, the data is already compromised.

But the technical limitations of spreadsheets are only part of the problem. The deeper issue is cognitive. Every time a PR team sets up a system to classify coverage, they face the same questions:

- What categories should we use?
- How detailed should they be?
- Should we tag by topic, message, format, spokesperson, or campaign?
- What's the difference between thought leadership and executive commentary?
- How do we keep this consistent over time?

These are not trivial decisions. They require judgement, experience, and constant maintenance. And yet, teams are expected to make them from scratch, every time.

The result is enormous cognitive load. Time and energy are spent designing classification systems rather than analysing the coverage itself. Many teams never get past this stage. They stall, simplify, or abandon structure altogether.

Insight is delayed not because teams lack ambition, but because they are stuck deciding how to organise the data before they can even begin asking questions of it.

In response to this complexity, free tagging offered in some software solutions appears to offer an escape.

At first glance, it sounds ideal. Let people tag coverage however they like. Let structure emerge organically. In practice it doesn't work.

Free tagging introduces:

- No shared definitions
- No consistency across teams
- No reliable comparability
- No ability to scale
- No dependable foundation for AI

Tags drift in meaning. Synonyms proliferate. Slight differences become permanent fractures in the data. Over time, the system becomes harder to use, not easier. Free tagging doesn't reduce effort, it pushes the burden of interpretation downstream, where it becomes harder and more expensive to resolve.

Most importantly, free tagging destroys trust. When no one is confident that tags mean the same thing everywhere, insight becomes subjective. Analysis turns into debate. Decisions slow down.

And when AI is introduced into this environment, the problems compound. AI cannot resolve ambiguity it wasn't designed to understand. Without consistent structure, it amplifies noise and produces answers that feel authoritative but rest on unstable ground. This is why free tagging fails, not in theory, but in practice.

We don't believe free tagging enables insight. We believe it guarantees chaos. Structure does not restrict PRs, it liberates and enables them.

A guided taxonomy removes cognitive load. It gives teams a shared language. It creates consistency without constant negotiation. It allows data to accumulate value over time rather than decay.

When structure is designed well, it doesn't limit what teams can learn. It makes learning possible in the first place. It's about creating a foundation strong enough to support analysis, comparison, automation, and AI-assisted insight.

The PR industry has spent years trying to extract intelligence from systems that were never designed to produce it. Spreadsheets and free tagging have taken us as far as they can. If PR is to deliver on the expectations placed on it today, real-time insight, market intelligence, and AI-ready data, it needs a different approach. It needs shared structure.

And that leads to the next question. If free tagging and spreadsheets are not the answer, what does a practical, industry-ready structure actually look like?

## Part 3 - Structure a prerequisite for insight

Insight does not emerge from data by accident. It emerges when data is consistent, comparable, and reliable enough to be questioned with confidence. Without those conditions, analysis becomes interpretation, and interpretation becomes opinion. This is as true in PR as it is in any other discipline.

For years, the PR industry has tried to extract insight from coverage that is inconsistently labelled, loosely categorised, and stored across fragile systems. When results fall short, the assumption is often that PR is inherently difficult to measure. But insight doesn't fail because PR is unmeasurable. Insight fails because the data isn't structured.

**At its core, insight is about relationships in data.**

To understand what is working, you need to compare like with like. To identify trends, you need consistency over time. To spot correlations, you need stable categories that mean the same thing from one dataset to the next.

Without structure comparisons break down, trends become unreliable, correlations are impossible to trust and analysis becomes anecdotal.

When the same type of coverage is labelled differently by different people, or across different periods, the data cannot be queried with confidence. Any patterns that appear are fragile at best, misleading at worst. This is why teams often feel they are "looking at the data" but not learning from it. The problem is not effort or intent, it's that the data does not support interrogation.

PR teams are increasingly being asked questions that go beyond visibility:

- Which narratives are gaining traction?
- Which messages are landing in which markets?
- What type of coverage drives search interest or trust?
- Where are early signals of reputational risk emerging?

These questions cannot be answered without consistent classification.

If a message category changes definition halfway through the year, year-on-year comparison becomes meaningless. If leadership coverage is tagged differently across teams or regions, market-level analysis collapses. If formats, topics, or spokespersons are inconsistently labelled, correlation with business signals becomes guesswork.

Causation and correlation are not advanced analytics problems. They are basic data hygiene problems. Without a stable structure, even the most sophisticated analysis will fail. AI has intensified this issue rather than solving it.

AI systems are exceptionally good at spotting patterns, summarising large volumes of content and answering questions based on available data, but they rely on one critical condition: structured inputs.

When AI is pointed at clean, consistently classified data, it becomes a powerful accelerator. It can surface trends quickly, highlight anomalies, and suggest lines of inquiry that would be time-consuming to uncover manually.

When AI is pointed at messy data, the opposite happens. It produces hallucinations, false narratives, misleading “insights”, confident answers built on weak foundations. This is not a flaw in AI. It is a predictable outcome of poor structure.

AI does not know when data is inconsistent unless it is explicitly told. It will happily draw connections between categories that should never have been compared, and it will present those connections with confidence. It is the difference between AI being helpful and AI being harmful.

To support reliable insight (human or machine) PR needs a single point of truth. Not a monitoring feed full of noise. Not a collection of disconnected spreadsheets. But a curated dataset that teams can trust.

A true single point of truth is:

- curated — coverage is intentionally selected, not indiscriminately collected
- human-selected — relevance is decided by people, not algorithms alone
- classified to a standard — categories mean the same thing everywhere
- junk-free — no duplication, no irrelevant clutter
- trustworthy — data can be queried without second-guessing its validity

This kind of dataset accumulates value over time. Each new piece of coverage strengthens the system rather than destabilising it. And once it exists, something important changes. Questions suddenly become answerable.

With structured, trustworthy data, PR teams can stop asking whether insight is possible and start asking better questions.

Questions like:

- How visible is our CEO in each market, and how has that changed over time?
- Which products are gaining momentum, and where?
- Which key messages consistently appear in high-quality coverage?
- What types of coverage actually correlate with search demand or trust indicators?
- Which narratives are emerging before they show up in formal reporting?
- Are there early signals of reputational risk we should be paying attention to?
- How effective is our PR investment across different strategies or markets?

Practical, operational, and time-sensitive questions that are only answerable when the underlying data is structured well enough to be queried with confidence.

Structure is what makes intelligence possible, without it, PR remains stuck in explanation and justification. With it, PR becomes a source of market insight.

That brings us to the next step. If structure is the prerequisite for insight, what does a practical, industry-ready structure actually look like?



## Part 4 - A coverage taxonomy standard

PR cannot deliver consistent insight without structure, AI cannot be trusted without clean inputs and asking every team to solve this problem independently has failed. The question, then, is not whether PR needs a standard coverage data taxonomy, but why the industry has gone so long without one.

It turns out that designing a usable taxonomy is deceptively difficult. On the surface, it looks like a simple categorisation exercise. In reality, it requires balancing competing needs: precision without complexity, flexibility without chaos, and standardisation without losing relevance across different organisations, markets, and campaign types.

Despite the complexity involved, the industry has quietly normalised something unreasonable. Every PR team is expected to:

- design its own classification system
- define its own categories
- decide its own naming conventions
- maintain consistency over time
- retrain new joiners
- rebuild the system when it breaks

All while delivering campaigns, managing stakeholders, responding to crises, and producing insight. This is not empowering. It is inefficient.

Worse, it guarantees fragmentation. Even when teams are thoughtful and capable, the resulting systems cannot be compared, shared, or scaled. Each taxonomy becomes a dead end, useful only to the people who created it, for a limited period of time. No other mature discipline works this way.

In other parts of the organisation, structure is assumed. CRM systems don't ask teams to invent lifecycle stages. Analytics platforms don't start with blank dashboards. Finance systems don't allow free-form definitions of revenue. Product teams don't track events without schemas.

These systems work because core structures are predefined, categories are shared, definitions are stable and learning accumulates over time. This doesn't limit teams. It enables them. By removing the burden of designing the system, teams can focus on using it, asking better questions, spotting patterns, and making decisions faster. PR deserves the same treatment. PR deserves infrastructure that reflects its importance, complexity, and growing strategic role. Not improvised systems built from scratch in spreadsheets.

A shared taxonomy does not magically produce insight on its own. But it creates the conditions in which insight becomes possible, at speed, at scale, and with confidence. It removes cognitive load. It eliminates reinvention. It restores trust in the data.

And it allows PR to finally behave like a data discipline, empowered to explore what becomes possible once that foundation is in place. When coverage data is no longer something to wrestle with, but something that works for you.

## Part 5 - Retrospective reporting to real-time intelligence

For most of its history, PR reporting has been retrospective. Coverage is gathered, summarised, and presented after the fact, often at the end of a month or campaign. Reports arrive as static PDFs or slide decks, designed to document what happened rather than inform what should happen next.

That model made sense when PR moved at a slower pace and reporting was primarily about justification. But it no longer matches how organisations operate, or what they now expect from communications teams. The future of PR measurement is not better reports. It is real-time intelligence.

Traditional PR reports answer a narrow set of questions:

- How much coverage did we get?
- Where did it appear?
- What was the reach?

Even when those reports are well-produced, they tend to have a short shelf life. By the time they are read, shared, and discussed, the moment they describe has already passed.

More importantly, they are rarely designed to be interrogated. Once a report is published, the structure is fixed. If a new question arises about a specific market, leader, narrative, or risk, the analysis often has to start again from scratch. This creates friction at exactly the wrong moment. Modern organisations need insight at the speed of decision-making, not on a reporting schedule.

When coverage data is consistently classified and stored as a single point of truth, reporting no longer has to be a one-off event. It becomes a live system. Instead of producing one report for everyone, PR teams can surface insight continuously, sliced and filtered to reflect the questions different stakeholders are asking.

Insight can be viewed by audience, by market, by product or theme, by leadership visibility, by risk or issue, over time, not just in snapshots.

Reporting becomes flexible, exploratory, and responsive. Weekly summaries, daily check-ins, and ad hoc analysis all draw from the same underlying data. Nothing needs to be rebuilt. Nothing needs to be reclassified. The answers are already there.

AI plays a critical role in this shift, not as a replacement for human judgement, but as an accelerator of discovery.

When applied to structured PR data, AI can:

- suggest interesting questions to ask
- surface anomalies or unexpected changes
- highlight emerging narratives before they become obvious
- identify patterns across markets or time periods
- summarise complex datasets quickly and clearly

This changes the nature of reporting.

Instead of deciding in advance what to include in a report, teams can explore the data in response to what is happening now. AI helps guide attention, pointing to what may matter, while humans decide what to do with that insight. The result is not automation for its own sake, but a shift from preparing answers to exploring signals.

As PR data becomes more intelligent, it also becomes more contextual. Not everyone needs to see the same thing.

- Leadership teams want high-level narrative and reputational signals.
- Crisis teams need early warnings and issue-specific views.
- Product teams want to understand momentum and messaging.
- Communications teams need operational detail and performance insight.
- Agencies need clarity, comparability, and shared definitions.

A real-time intelligence model allows each of these groups to access the same underlying data, but through different lenses.

This is not fragmentation. It is alignment. Everyone is working from the same source of truth, but seeing the slice that matters to them.

The most important change is not technical. It is conceptual. PR stops reporting what happened and starts explaining what's happening. That shift turns PR from a function that looks backwards into one that helps organisations understand the present and prepare for what comes next.

It also changes how PR teams are perceived internally, from reporters of activity to interpreters of the market. This is what becomes possible when structure, consistency, and intelligence are built into the system from the start..

## Part 6 - From structure to practice

This report has focused on the problem PR faces and the structural gap that sits beneath it. It has argued for the importance of shared classification, consistent definitions, and a common foundation for coverage data. But structure only matters if it works in practice.

No taxonomy proves its value in workshops or documents. It proves itself when teams are under pressure, deadlines are tight, stakeholders are asking questions, and reporting needs to be done quickly and accurately. That is where structure either holds up or falls apart.

Consistency is not achieved through discussion. It is achieved through repetition. Edge cases do not appear in planning sessions. They appear in day-to-day workflows, where real coverage does not fit neatly into theoretical categories. That is why this work is happening in practice, not in theory.

We are already working with thousands of PR teams all over the world already using CoverageBook. This provides a live environment where structure can be tested against real reporting needs, across different sectors, markets and team sizes.

Rather than introducing an entire taxonomy at once, we are rolling this out deliberately, category by category, with care. Definitions will be explained clearly, examples provided, the purpose of each category made explicit, not assumed. This approach avoids overwhelming teams with complexity before value is clear. It allows structure to be learned, applied, and understood as part of normal reporting, rather than as an additional burden.

We are starting with the categories that matter most in day-to-day PR reporting such as coverage type, key messaging, and tier or quality indicators. These are the categories that appear in almost every report, yet are most commonly defined differently from one team to the next, making comparison and insight difficult. By starting here, we are focusing on the foundations that unlock the most immediate value. As each category is rolled out, questions surface naturally. Ambiguities become visible. Edge cases appear. Patterns emerge.

Refinement will be driven by what teams actually struggle with, not by abstract preference. Changes are made where definitions cause confusion, where categories break down, or where structure creates unnecessary friction. This is how the taxonomy will evolve: through evidence, not opinion.

We are not waiting for perfect consensus. We are not freezing this as a fixed system. And we are not treating this as a one-off release. We are building something durable, improving it through use, and letting usefulness determine adoption.

Coverage data has outgrown improvised systems. PR reporting now demands structure that can scale, support insight, and withstand the growing role of automation and AI. This is our response to that reality.

We are building it in the open, refining it in practice, and sharing what we learn as we go. If the themes in this report resonate and you want to be part of that journey, join us and thousands of PR practitioners at CoverageBook, where this PR data revolution will take place.

# Context

The ideas in this book are shaped by decades of experience working at the intersection of PR, search, and measurement, where structure, classification, and standards determine what can be understood and what gets lost.

Before focusing exclusively on PR measurement, Gary Preston and Jack Hubbard spent many years working on how information is discovered, interpreted, and trusted at scale. That work included applying PR principles to search marketing well before those disciplines formally converged.

Over time, this led to the creation of tools designed to make complex data more accessible to practitioners, helping teams understand demand, narrative momentum, and visibility without adding unnecessary complexity. Those platforms reached global audiences before attention shifted fully toward PR measurement and evaluation.

More recent work has focused on connecting coverage to business outcomes and understanding how large language models reference media sources when forming responses, reinforcing the importance of structure, curation, and consistency in modern communications data.

Across all of this work, the common thread has been the same:

- a belief that useful tools are built by solving real problems
- that clarity matters more than novelty
- that long-term value comes from getting the foundations right

That perspective informs everything in this report.